



NANODEGREE PROGRAM SYLLABUS

Artificial Intelligence



Overview

This program will teach you how to become a better Artificial Intelligence or Machine Learning Engineer by teaching you classical AI algorithms applied to common problem types. You will complete projects and exercises incorporating search, optimization, planning, and probabilistic graphical models which have been used in Artificial Intelligence applications for automation, logistics, operations research, and more. These concepts form the foundation for many of the most exciting advances in AI in recent years. Each project you build will be an opportunity to demonstrate what you've learned in your lessons, and become part of a career portfolio that will demonstrate your mastery of these skills to potential employers.

In order to succeed in this program, we recommend having experience with intermediate Python programming (including experience with basic algorithms, common data structures, and Object Oriented Programming), and intermediate statistics & linear algebra (including discrete & continuous distributions, vector spaces & matrices). If you'd like to prepare for this Nanodegree program, start with our [AI Programming with Python](#) Nanodegree program, then complete either our [Machine Learning Engineer](#) or [Deep Learning](#) Nanodegree programs.


This program covers classical AI techniques that you will need to master to become a better AI practitioner. Specifically, we will focus on intermediate to advanced programming skills, linear algebra, and algorithms that appear in a variety of AI applications.



Estimated Time:
3 Months at
12-15 hours / week



Prerequisites:
Algebra, Calculus,
Statistics, &
Python



Flexible Learning:
Self-paced, so
you can learn on
the schedule that
works best for you



Need Help?
[udacity.com/advisor](https://www.udacity.com/advisor)
Discuss this program
with an enrollment
advisor.

Course 1: Intro to Artificial Intelligence

LEARNING OUTCOMES

LESSON ONE

Welcome to the Program

- Meet the course instructors and Udacity team. Learn about the resources available to help you succeed

LESSON TWO

Intro to Artificial Intelligence

- Consider the meaning of “artificial intelligence”
- Be able to define core concepts from AI including “agents”, “environments”, and “states”
- Learn the concept of “rational” behavior for AI agents

LESSON THREE

Setting Up your Environment with Anaconda

- Install the software and complete necessary system configuration you’ll need for the projects

Course 2: Constraint Satisfaction Problems

Use constraint propagation and search to build an agent that reasons like a human would to efficiently solve any Sudoku puzzle.

Course Project: Build a Sudoku Solver

Humans use reason to solve problems by decomposing the problem statement and incorporating domain knowledge to limit the possible solution space. In this project you'll use a technique called constraint propagation together with backtracking search to make an agent that only considers reasonable solution candidates and efficiently solves any Sudoku puzzle. This approach appears in many classical AI problems, and the solution techniques have been extended and applied to diverse problems in bioinformatics, logistics, and operations research.

In this project you will demonstrate some basic algorithms knowledge, and learn to use constraint satisfaction to solve general problems.

LEARNING OUTCOMES

LESSON ONE

Solving Sudoku With AI

- Express logical constraints as Python functions
- Use constraint propagation & search to solve all Sudoku puzzles

LESSON TWO

Constraint Satisfaction Problems

- Learn to represent problems in terms of logical constraints
- Use constraint propagation to limit the potential solution space
- Incorporate backtracking search to find a solution when the set of constraints is incomplete

LESSON THREE

Additional Topics in CSP

- List of external resources for you to continue learning about CSPs

Course 3: Search, Optimization, and Planning

Build agents that can reason to achieve their goals using search and symbolic logic—like the NASA Mars rovers.

Course Project: Build a Forward Planning Agent

Intelligent agents are expected to act in complex domains where their goals and objectives may not be immediately achievable. They must reason about their goals and make rational choices of actions to achieve them. In this project you will build a system using symbolic logic to represent general problem domains and use classical search to find optimal plans for achieving your agent's goals. Planning & scheduling systems power modern automation & logistics operations, and aerospace applications like the Hubble telescope & NASA Mars rovers.

In this project you will demonstrate an understanding of classical optimization & search algorithms, symbolic logic, and domain-independent planning.

LEARNING OUTCOMES

LESSON ONE

Introduction

- Learn about the significance of search in AI

LESSON TWO

Uninformed Search

- Learn uninformed search techniques including depth-first order, breadth-first order, and Uniform Cost Search

LESSON THREE

Informed Search

- Learn informed search techniques (using heuristics) including A*
- Understand admissibility and consistency conditions for heuristics

LESSON FOUR

Additional Topics: Search

- List of external resources for you to continue learning about search

LESSON FIVE**Classroom Exercise:
Search**

- Implement informed & uninformed search for Pacman

LESSON SIX**Introduction:
Optimization
Problems**

- Introduce iterative improvement problems that can be solved with optimization

LESSON SEVEN**Hill Climbing**

- Learn Random Hill Climbing for local search optimization problems

LESSON EIGHT**Simulated Annealing**

- Learn to use Simulated Annealing for global optimization problems

LESSON NINE**Genetic Algorithms**

- Explore and implement Genetic Algorithms that keep a pool of candidates to solve optimization problems

LESSON TEN**Additional
Optimization
Topics**

- Learn about improvements & optimizations to optimization search including Late Acceptance Hill Climbing, Basin Hopping, & Differential Evolution

LESSON ELEVEN**Classroom Exercise:
Optimization
Problems**

- Compare optimization techniques on a variety of problems

LESSON TWELVE**Symbolic Logic &
Reasoning**

- Learn Propositional logic (propositions & statements)
- Learn First-Order logic (quantifiers, variables, & objects)
- Encode problems with symbolic constraints using first-order logic

LESSON THIRTEEN**Introduction to
Automated
Planning**

- Learn to define planning problems

LESSON FOURTEEN

Classical Planning

- Learn high-level features of automated planning techniques using search & symbolic logic including forward planning, backwards planning, & hierarchical planning
- Explore planning heuristics & planning graphs

LESSON FIFTEEN

Additional Topics in Planning

- List of external resources for you to continue learning about planning



Course 4: Adversarial Search

Extend classical search to adversarial domains, to build agents that make good decisions without any human intervention—such as the DeepMind AlphaGo agent.

Course Project: Build an Adversarial Game Playing Agent

Intelligent agents are expected to act in complex domains where their goals and objectives may not be immediately achievable. They must reason about their goals and make rational choices of actions to achieve them. In this project you will build a system using symbolic logic to represent general problem domains and use classical search to find optimal plans for achieving your agent's goals. Planning & scheduling systems power modern automation & logistics operations, and aerospace applications like the Hubble telescope & NASA Mars rovers.

In this project you will demonstrate an understanding of classical optimization & search algorithms, symbolic logic, and domain-independent planning.

LEARNING OUTCOMES

LESSON ONE

Search in Multi-Agent Domains

- Understand “adversarial” problems & applications (e.g., multi-agent environments)
- Extend state space search techniques to domains your agents do not fully control
- Learn the minimax search technique

LESSON TWO

Optimizing Minimax Search

- Learn techniques used to overcome limitations in basic minimax search like depth-limiting and alpha-beta pruning,

LESSON THREE

Extending Minimax Search

- Extend adversarial search to non-deterministic domains and domains with more than two players

LESSON FOUR

Additional Adversarial Search Topics

- List of external resources for you to continue learning about adversarial search

Course 5: Fundamentals of Probabilistic Graphical Models

Model real-world uncertainty through probability to perform pattern recognition.

Course Project: Part of Speech Tagging

Probabilistic models allow your agents to better handle the uncertainty of the real world by explicitly modeling their belief state as a distribution over all possible states. In this project you'll use a Hidden Markov Model (HMM) to perform part of speech tagging, a common pre-processing step in Natural Language Processing. HMMs have been used extensively in NLP, speech recognition, bioinformatics, and computer vision tasks.

LEARNING OUTCOMES

LESSON ONE

Probability

- Review key concepts in probability including discrete distributions, joint probabilities, and conditional probabilities

LESSON TWO

Bayes Networks

- Efficiently encode joint probabilities in Bayes networks

LESSON THREE

Inference in Bayes Nets

- Learn about inference in Bayes networks through exact enumeration with optimizations
- Learn techniques for approximate inference in more complex Bayes networks

LESSON FOUR

Hidden Markov Models

- Learn parameters to maximize the likelihood of model parameters to training data
- Determine the likelihood of observing test data given a fixed model
- Learn an algorithm to Identify the most likely sequence of states in a model given some data

LESSON FIVE

Dynamic Time Warping

- Learn the dynamic time warping algorithm for time-series analysis


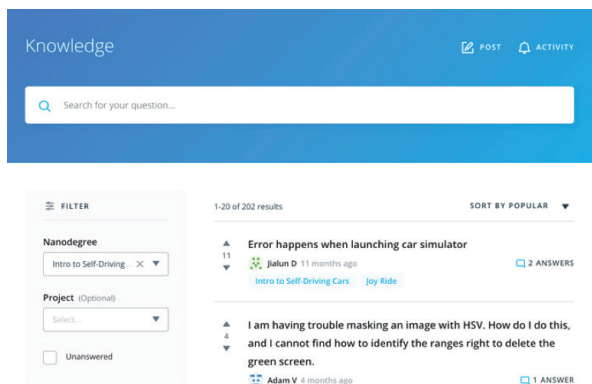

LESSON SIX

Additional Topics in PGMs

- List of external resources for you to continue learning about probabilistic graphical models



Our Classroom Experience



REAL-WORLD PROJECTS

Build your skills through industry-relevant projects. Get personalized feedback from our network of 900+ project reviewers. Our simple interface makes it easy to submit your projects as often as you need and receive unlimited feedback on your work.

KNOWLEDGE

Find answers to your questions with Knowledge, our proprietary wiki. Search questions asked by other students, connect with technical mentors, and discover in real-time how to solve the challenges that you encounter.

STUDENT HUB

Leverage the power of community through a simple, yet powerful chat interface built within the classroom. Use Student Hub to connect with your fellow students in your Executive Program.

WORKSPACES

See your code in action. Check the output and quality of your code by running them on workspaces that are a part of our classroom.

QUIZZES

Check your understanding of concepts learned in the program by answering simple and auto-graded quizzes. Easily go back to the lessons to brush up on concepts anytime you get an answer wrong.

CUSTOM STUDY PLANS

Preschedule your study times and save them to your personal calendar to create a custom study plan. Program regular reminders to keep track of your progress toward your goals and completion of your program.

PROGRESS TRACKER

Stay on track to complete your Nanodegree program with useful milestone reminders.

Learn with the Best



Peter Norvig

RESEARCH DIRECTOR

Peter Norvig is a Director of Research at Google and is co-author of *Artificial Intelligence: A Modern Approach*, the leading textbook in the field.



Sebastian Thrun

INSTRUCTOR

As the founder and president of Udacity, Sebastian's mission is to democratize education. He is also the founder of Google X, where he led projects including the Self-Driving Car, Google Glass, and more.



Thad Starner

PROFESSOR OF COMPUTER SCIENCE

Thad Starner is the director of the Contextual Computing Group (CCG) at Georgia Tech and is also the longest-serving Technical Lead/Manager on Google's Glass project.

All Our Nanodegree Programs Include:



EXPERIENCED PROJECT REVIEWERS

REVIEWER SERVICES

- Personalized feedback & line by line code reviews
- 1600+ Reviewers with a 4.85/5 average rating
- 3 hour average project review turnaround time
- Unlimited submissions and feedback loops
- Practical tips and industry best practices
- Additional suggested resources to improve



TECHNICAL MENTOR SUPPORT

MENTORSHIP SERVICES

- Questions answered quickly by our team of technical mentors
- 1000+ Mentors with a 4.7/5 average rating
- Support for all your technical questions



PERSONAL CAREER SERVICES

CAREER COACHING

- Personal assistance in your job search
- Monthly 1-on-1 calls
- Personalized feedback and career guidance
- Interview preparation
- Resume services
- Github portfolio review
- LinkedIn profile optimization



Frequently Asked Questions

PROGRAM OVERVIEW

WHY SHOULD I ENROLL?

The Artificial Intelligence Nanodegree program features expert instructors, and world-class curriculum built in collaboration with top companies in the field. The program offers a broad introduction to the field of artificial intelligence, and can help you maximize your potential as an artificial intelligence or machine learning engineer. If you're ready for an efficient and effective immersion in the world of AI, with the goal of pursuing new opportunities in the field, this is an excellent program for you.

WHAT JOBS WILL THIS PROGRAM PREPARE ME FOR?

This Nanodegree program is designed to build on your existing skills as an engineer or developer. As such, it doesn't prepare you for a specific job, but instead expands your skills with artificial intelligence algorithms. These skills can be applied to various applications such as video game AI, pathfinding for robots, and recognizing patterns over time like handwriting and sign language.

HOW DO I KNOW IF THIS PROGRAM IS RIGHT FOR ME?

In this Nanodegree program, you will learn from the world's foremost AI experts, and develop a deep understanding of algorithms being applied to real-world problems in Natural Language Processing, Computer Vision, Bioinformatics and more. If your goal is to become an AI expert, then this program is ideal, because it teaches you some of the most important algorithms in AI. You'll benefit from a structured approach for applying these techniques to new challenges, and emerge from the program fully prepared to advance in the field.

ENROLLMENT AND ADMISSION

DO I NEED TO APPLY? WHAT ARE THE ADMISSION CRITERIA?

There is no application. This Nanodegree program accepts everyone, regardless of experience and specific background.

WHAT ARE THE PREREQUISITES FOR ENROLLMENT?

If you are new to Python programming, we suggest our [AI Programming with Python Nanodegree](#) program. You must have completed a course in computer science algorithms equivalent to the [Data Structures & Algorithms Nanodegree](#) program prior to entering the program. Additionally, you should have the following knowledge:

Intermediate Python programming knowledge, including:

- Strings, numbers, and variables
- Statements, operators, and expressions
- Lists, tuples, and dictionaries
- Conditions & loops



FAQs Continued

- Generators & comprehensions
- Procedures, objects, modules, and libraries
- Troubleshooting and debugging
- Research & documentation
- Problem solving
- Algorithms and data structures

Basic shell scripting:

- Run programs from a command line
- Debug error messages and feedback
- Set environment variables
- Establish remote connections

Basic statistical knowledge, including:

- Populations, samples
- Mean, median, mode
- Standard error
- Variation, standard deviations
- Normal distribution

Intermediate differential calculus and linear algebra, including:

- Derivatives & Integrals
- Series expansions
- Matrix operations through eigenvectors and eigenvalues

Additionally, you should be able to follow and interpret pseudocode for algorithms like the example below and implement them in Python. You should also be able to informally evaluate the time or space complexity of an algorithm. For example, you should be able to explain that a for loop that does constant $O(1)$ work on each iteration over an array of length n has a complexity of $O(n)$.

```
function Hill-Climbing(problem) returns a State
  current <- Make-Node(problem.Initial-State)
  loop do
    neighbor <- a highest-valued successor of current
    if neighbor.value ≤ current.value then return current.state
  current <- neighbor
```

IF I DO NOT MEET THE REQUIREMENTS TO ENROLL, WHAT SHOULD I DO?

We have a number of courses and programs we can recommend that will help prepare you for the program, depending on the areas you need to address. For example:

- [Intro to Machine Learning](#)
- [Artificial Intelligence Programming with Python Nanodegree program](#)
- [Data Structures & Algorithms Nanodegree program](#)
- [Data Analyst Nanodegree program](#)
- [Intro to Machine Learning](#)
- [Machine Learning Engineer Nanodegree program](#)



FAQs Continued

TUITION AND TERM OF PROGRAM

HOW IS THIS NANODEGREE PROGRAM STRUCTURED?

The Artificial Intelligence Nanodegree program is comprised of content and curriculum to support four (4) projects. We estimate that students can complete the program in three (3) months working 10 hours per week.

Each project will be reviewed by the Udacity reviewer network. Feedback will be provided and if you do not pass the project, you will be asked to resubmit the project until it passes.

HOW LONG IS THIS NANODEGREE PROGRAM?

Access to this Nanodegree program runs for the length of time specified in the payment card above. If you do not graduate within that time period, you will continue learning with month to month payments. See the [Terms of Use](#) and [FAQs](#) for other policies regarding the terms of access to our Nanodegree programs.

CAN I SWITCH MY START DATE? CAN I GET A REFUND?

Please see the Udacity Nanodegree program [FAQs](#) for policies on enrollment in our programs.

SOFTWARE AND HARDWARE

WHAT SOFTWARE AND VERSIONS WILL I NEED IN THIS PROGRAM?

You will need a computer running a 64-bit operating system (most modern Windows, OS X, and Linux versions will work) with at least 8GB of RAM, along with administrator account permissions sufficient to install programs including Anaconda with Python 3.6 and supporting packages. Your network should allow secure connections to remote hosts (like SSH). We will provide you with instructions to install the required software packages.

